

Prudentia AI Administrator Manual

Last updated: 2026-07-07

This manual explains how to administer Prudentia AI from the web application. It is written for Platform Admins, System Admins, and Space Admins. For startup, backup, update, and Docker runtime work, use the Operator Manual (operator-manual.md).

1. Administrator Roles

Role	Main responsibility	Important limits
Platform Admin	Full platform ownership, Runtime Settings, RAG Evaluation, users, spaces, audit, documents, connectors	Use sparingly. This role can change runtime behavior for the whole deployment.
System Admin	Global operations, users except Platform Admins, spaces, documents, review, audit, recovery	No Platform Admin-only Runtime Settings or RAG Evaluation access.
Space Admin	Knowledge Space governance, document work, review, and scoped user management	Can manage allowed spaces, documents, connectors, and lower-scope users, but not global platform settings or peer/admin accounts.

Account Type Assignment Reference

Use the narrowest role that lets the person do their job.

Account type	Why it exists	Use it when the person needs to	Do not use it when
Platform Admin	Full platform ownership	Manage every user and space, change Runtime Settings, run evaluations, govern connectors, review audit, and make deployment-wide model choices	The person only needs operational or space-level administration
System Admin	Global operations without Runtime Settings	Manage users except Platform Admins, govern documents, recover ingestion, review audit, resolve Review Queue holds, and manage spaces globally	The person must change model/provider settings or run RAG Evaluation
Space Admin	Space-level administration	Manage assigned spaces, scoped Chat Member and Document Contributor accounts, uploads, folder sources, connector scopes, document lifecycle, and Review Queue work	The person needs global configuration, RAG Evaluation, audit review, or management of global/peer admin accounts
Document Contributor	Document intake, review, and maintenance	Query, upload, track, resolve Review Queue items, reingest, restore, or trash writable documents in assigned spaces	The person should only ask questions and inspect evidence, or needs user/space administration

Account type	Why it exists	Use it when the person needs to	Do not use it when
Chat Member	Everyday knowledge use	Ask grounded questions, inspect citations, and browse visible document metadata	The person needs to upload, manage, review, audit, or configure anything
Audit Viewer	Governance review	Search, filter, inspect, and export visible audit events without changing system state	The person needs document content access through Query Intelligence

User Manager and Reviewer are legacy account types retained for old records and compatibility. New assignments should use **Space Admin** for scoped user management and **Document Contributor** for upload plus review.

If a page is not visible, the signed-in administrator does not have the role required for that page.

2. First Platform Admin Checklist

After the operator starts the deployment:

1. Open the app URL. Local deployments commonly use `http://localhost:3000`.
2. Sign in with `BOOTSTRAP_ADMIN_EMAIL` and `BOOTSTRAP_ADMIN_PASSWORD`.
3. Open **Account Management** and change the bootstrap password immediately.
4. Open **Govern > Runtime Settings** and confirm runtime model health.
5. Confirm the active provider stack, role assignments, query planner, and ingestion profile.
6. Create the top-level Knowledge Spaces.
7. Create System Admin, Space Admin, Document Contributor, Chat Member, and Audit Viewer accounts as needed.
8. Upload a small pilot document and confirm it reaches **Indexed** status.
9. Ask a cited smoke-test question in **Query Intelligence**.
10. If Live DB is required, create and approve the first least-privilege database scope.
11. Review **System Audit** for the expected bootstrap, login, user, document, and configuration events.

3. User Management

Open **Govern > User Management**.

Platform Admins and System Admins manage users globally according to their role limits. Space Admins manage only Chat Member and Document Contributor accounts inside their assigned Knowledge Spaces and clearance scope.

Create A User

1. Select **Create user**.
2. Enter email and display name.
3. Choose an account type you are allowed to assign.

4. Choose a clearance level. Global admin roles receive the maximum effective clearance.
5. Select one or more Knowledge Space memberships.
6. Keep **Active account** selected unless the user should not sign in yet.
7. Copy or regenerate the initial password.
8. Select **Create user**.
9. Hand off the password securely. The user can change it after signing in.

Copy the initial password before closing the dialog. It is the handoff credential.

Edit A User

1. Select the edit action on the user row.
2. Update name, account type, clearance, active state, or Knowledge Spaces.
3. Save the user.

If you edit your own account, your session context may refresh after saving.

Reset A Password

Use this when a user forgets their password or a temporary credential must be issued.

1. Select the reset-password action on the user row.
2. Copy or regenerate the temporary password.
3. Select **Reset password**.
4. Copy the temporary password before closing the dialog.
5. Give the temporary password to the user through an approved secure channel.

Existing sessions for that user are signed out. The user must change the temporary password at next sign-in. You cannot reset your own password from User Management; use **Account Management** for your own password.

Disable Or Delete A User

Disable an account when access should be blocked but memberships should remain available for later review. Delete an account only when policy allows removal.

Deleting a user removes the account, memberships, chat history, and generated artifacts. Documents and audit history are retained. User deletion requires typing the target email address to confirm.

Chat Activity

Platform Admins can inspect saved chat activity for managed users when the action appears in User Management. Use this for governance and investigation only, according to organizational policy.

4. Knowledge Space Management

Knowledge Spaces define document and live-data access boundaries. Space names use path-like values such as `/finance`, `/legal/contracts`, or `/manuals`.

Space Admins and global admins can create, edit, and delete allowed spaces. Deleting a Knowledge Space only works when the space is empty, has no child spaces, and has no assigned users.

Recommended practice:

- Create top-level spaces before onboarding users.
- Assign users only to the spaces they need.
- Use child spaces when access should be narrower than the parent department.
- Keep space names stable because users, documents, folder sources, and database scopes depend on them.

5. Clearance Levels

Documents and users can carry these clearance levels:

- NATO_UNCLASSIFIED
- NATO_RESTRICTED
- NATO_CONFIDENTIAL
- NATO_SECRET
- COSMIC_TOP_SECRET

Users can access documents at or below their effective clearance level inside their allowed Knowledge Spaces. Do not assign higher clearance than the user is authorized to hold.

6. Document Governance

Use **Corpus > Document Library** to inspect and govern documents.

Common actions:

Action	Effect
Download	Opens or downloads the original source file.
Reingest	Queues a new ingestion job from the stored source file.
Move to Trash	Soft deletes the document and removes indexed vectors.
Restore	Restores a Trash document and queues reingestion.
Permanently Delete	Removes source file, image assets, vectors, and the document record where policy and role allow.

Use permanent deletion only when retention policy allows it. Prefer Trash when an item should be hidden from retrieval but retained for review.

When replacing a document, use supersedes metadata so the previous version stays auditable.

7. Folder Sources

Open **Corpus > Document Intake > Folder Sources** to upload local folder snapshots for bulk ingestion.

Folder snapshots are one-time schedules. They do not continuously watch the local folder. To ingest later file edits, choose the folder again and create a new snapshot.

Use run details to review queued, skipped, cancelled, and failed files. Unsupported files are recorded as skipped.

8. Database Connectors

Platform Admins, System Admins, and Space Admins can prepare governed live SQL access from **Document Intake > Database Connectors**.

Supported source types:

- SQL Server
- PostgreSQL

Database connectors do not copy database rows into the document corpus. At question time, Prudentia generates a read-only SQL query, validates it against an approved schema catalog, applies result limits and a timeout, and uses the returned rows as answer evidence.

Connector Approval Workflow

1. Add a connection profile with a least-privilege read-only database account.
2. Save the encrypted profile.
3. Test the profile.
4. Select **Read Schema** to capture schema metadata.
5. Select **Prepare Review** to create a reviewable schema catalog.
6. Review table and column descriptions, synonyms, relationships, business rules, allowed fields, and sensitivity flags.
7. Assign the owner Knowledge Space, optional shared Knowledge Spaces, and clearance.
8. Select **Continue AI Enrichment** if enrichment has not finished.
9. Select **Save Review** while drafting or while waiting for review.
10. Select **Enable Live DB Access** only after human review is complete.

Only approved, non-sensitive tables, columns, and relationships can be queried. Write statements, wildcard column selection, unapproved joins, system schemas, and chained SQL statements are blocked.

9. Review Queue Administration

Document Contributors, Space Admins, System Admins, and Platform Admins can open **Review Queue**.

If many documents pause in **Needs review**:

1. Review the OCR and PDF image thresholds in **Runtime Settings > Ingestion Controls**.
2. Confirm the ingestion worker is healthy in **Ingestion Health**.

3. Have Document Contributors or admins resolve the oldest OCR blocks and PDF image batches first.
4. Re-check Activity after approvals or skip decisions to confirm ingestion continues.

10. RAG Evaluation

RAG Evaluation is Platform Admin-only. Use it to measure answer quality with repeatable datasets.

Dataset cases should include a question or query. Recommended fields include stable IDs, expected_answer, expected source documents, acceptable source pages, minimum source count, degradation policy, must-include text, must-not-include text, citation requirements, faithfulness threshold, and latency threshold.

Supported optional expectation fields include expected_source_docs, acceptable_source_pages, min_sources, must_include, must_not_include, must_cite_source or requires_citations, min_faithfulness_score, allow_degraded, and latency_threshold_ms.

Basic workflow:

1. Open **Evaluate > RAG Evaluation**.
2. Import a JSON or JSONL dataset.
3. Choose the dataset.
4. Optionally set Knowledge Space, document IDs, case IDs, or a numeric limit.
5. Select **Run evaluation**.
6. Review pass rate, retrieval, reranking, citation, faithfulness, fallback, and latency metrics.
7. Inspect failed cases and retry after fixing data, ingestion, runtime, or expectations.

Active runs can be cancelled. Failed, partial, or cancelled runs can be retried.

Answer-content checks first use deterministic literal matching. When the evaluation answer LLM verifier is enabled, a failed literal must-include check can be reviewed semantically by the configured model. The verdict is recorded in the case **Checks and diagnostic JSON** as `answer_content.llm_verifier`.

11. Runtime Settings

Runtime Settings is Platform Admin-only under **Govern > Runtime Settings**. It includes **Models & Roles**, **Inference Services**, and **Ingestion Controls**.

Models & Roles

Use the guided flow in order:

1. Choose a starting stack: **Ollama local** or **vLLM text stack**.
2. Check service discovery for language, vision, embedding, and reranker providers.
3. Assign runtime roles:

Runtime role	Purpose
Answer synthesis	Produces the final grounded response.
Reasoning	Supports planning, rewrites, dates, and complex retrieval.
Router	Checks ambiguous routes before retrieval or file generation.
Faithfulness	Checks whether generated answers are grounded.
Ingestion metadata	Generates summaries, topics, document types, and claims.
Vision	Performs image OCR, captions, and optional PDF layout repair.
Embeddings	Creates dense vectors.
Reranker	Reorders retrieved evidence before synthesis.

4. Tune budgets, timeouts, model thinking, and the query planner.

5. Select **Test draft**.

6. Select **Save & activate** only after validation succeeds.

Changing the embedding provider or embedding model requires reindexing existing documents so stored vectors match the active embedding model.

Inference Services

Use **Inference Services** to check provider availability and manage vLLM launch limits. Applying vLLM limits restarts only the selected vLLM service and interrupts requests using that service.

Do not switch the active provider to vLLM unless the vLLM services and model caches are actually available. In the default Windows airgapped rollout, Ollama is the expected provider until vLLM is deliberately enabled.

Ingestion Controls

Use **Ingestion Controls** to set worker concurrency, OCR review threshold, PDF image review threshold, vision layout repair, and graph-enrichment visibility for new ingestion jobs.

One worker-concurrency slot is recommended for the default 8 GB Docker environment. Increasing concurrency can improve throughput only if CPU, memory, OCR, model, and storage capacity can support it.

The PDF image review threshold pauses unusually large image-analysis batches in Review Queue before vision analysis runs. Set it to 0 only when policy allows the image-review gate to be skipped. Graph enrichment controls whether the completed-document graph enrichment action is visible; keep it off for faster bulk ingestion when corpus-level graph analysis is not needed.

12. System Audit

Open **Govern > System Audit** to review append-only activity visible to your account.

You can filter by:

- Search text.
- Category.

- Event type.
- Actor ID or email.
- Target type.
- Target ID.
- Knowledge Space.
- Date range.

Use **Export CSV** to preserve a filtered audit set. The exported file should be handled according to your organization's sensitive-data policy.

13. Admin Troubleshooting

Problem	Likely cause	Admin action
User cannot sign in	Wrong password, inactive account, forced password change	Reset password, confirm active state, ask user to change temporary password.
User cannot see a page	Role does not include the route	Confirm account type and assign the correct role only if justified.
User cannot see documents	Wrong space, low clearance, lifecycle filter	Check Knowledge Space memberships, clearance, and document lifecycle.
Upload stuck in queued	Worker unavailable or capacity saturated	Check Ingestion Health and ask operator to inspect worker logs if needed.
Many Needs review jobs	OCR threshold, PDF image threshold, or low-quality scans	Assign Document Contributors or admins to clear review work, or tune thresholds after policy review.
Live DB source missing	Catalog not approved or outside user's scope	Review connector catalog, Knowledge Space, clearance, and disabled state.
Query answers degrade	Runtime provider, retrieval, or source availability issue	Check Runtime Settings health, document indexing, and source mode.
GraphRAG unavailable	Graph worker or Neo4j unavailable, or docs not opted in	Check Ingestion Health and operator runtime status.

14. Admin Safety Rules

- Use least privilege for roles, Knowledge Spaces, database scopes, and clearance.
- Do not approve database catalogs without human review.
- Do not broaden access just to make a query work.
- Do not change model provider or embedding model without a reindex plan.
- Do not permanently delete documents unless retention policy allows it.
- Document operational incidents with timestamps, job IDs, user IDs, document IDs, and screenshots where appropriate.