

Prudentia AI End User Walkthrough

Search documents and live data, choose ingestion profiles, verify evidence, and use optional GraphRAG

Field	Details
Audience	Chat Members, Document Contributors, Audit Viewers, Space Admins, System Admins, and Platform Admins
Scope	Document and live-data queries, connectors, ingestion profiles, optional GraphRAG, evidence, activity, governance, and Runtime Settings
Not covered	Deployment internals, secret rotation, database administration, and developer operations
Updated	2026-07-07

Tip: Read the Quick Start first, then jump to the feature that appears in your sidebar.

How to Use This Guide

Your sidebar is permission-aware. If a feature described here does not appear, your account does not include that permission.

1. Sign in with your workspace email and password.
2. Check the active Knowledge Space in the sidebar header.
3. Open Query Intelligence and ask a focused question.
4. Inspect citations before relying on an answer.
5. Use Add Files, Activity, Review Queue, or Audit only when those pages appear in your sidebar.

Know Your Account Type

Your account type controls which pages and actions appear. Knowledge Space membership and clearance still limit what you can see inside those pages.

Account type	Typical access
Platform Admin	Full platform ownership: users, spaces, Runtime Settings, RAG Evaluation, audit, connectors, uploads, review, and documents.
System Admin	Global operations: users except Platform Admins, spaces, documents, uploads, review, audit, and recovery.
Space Admin	Assigned-space administration: scoped users, spaces, uploads, review, folder sources, connectors, and document lifecycle.
Document Contributor	Document work: chat, upload, track ingestion, resolve Review Queue items, and maintain writable documents.
Chat Member	Everyday use: chat with visible documents, scope questions, inspect evidence, and browse visible library metadata.
Audit Viewer	Governance review: inspect visible audit events and metadata without changing system state.

Legacy User Manager and Reviewer records may exist in older data. New assignments use Space Admin for scoped user management and Document Contributor for upload plus review.

Feature Map

Use this map to connect common tasks to the page where they happen.

Feature	What you do	Where to go
Query Intelligence	Ask questions, watch the retrieval progress, and inspect cited evidence.	Operate > Query Intelligence
Knowledge Space selector	Choose the document scope used for queries and uploads.	Sidebar header
Document Library	Browse, search, inspect, download, and manage visible documents.	Corpus > Document Library

Feature	What you do	Where to go
Add Files	Upload PDF, DOCX, JPG, JPEG, PNG, or JSON files and choose an ingestion profile.	Corpus > Document Intake > Add Files
Activity	Track upload, folder, restore, and reingestion jobs.	Corpus > Document Intake > Activity
Folder Sources	Upload local folder snapshots.	Corpus > Document Intake > Folder Sources
Database Connectors	Approve SQL Server or PostgreSQL scopes for live read-only answers.	Corpus > Document Intake > Database Connectors
Review Queue	Resolve OCR blocks and PDF image-analysis holds before indexing continues.	Evaluate > Review Queue
System Audit	Review visible authentication, document, and ingestion activity.	Govern > System Audit
Runtime Settings	Assign inference roles, tune services, and manage ingestion controls.	Govern > Runtime Settings

Sign In and Navigate

- Open the web application URL provided by your organization. Local deployments commonly use `http://localhost:3000`.
- Sign in with your email and password. Change the initial password in Account Management when needed.
- Use the sidebar sections to move through Operate, Corpus, Evaluate, and Govern pages.
- Use the account area in the sidebar footer to view session details, change password, switch theme, or sign out.

Work With Knowledge Spaces

Knowledge Spaces are the document boundaries used by the system.

- Select the active space before asking questions or uploading from chat.
- Queries retrieve from the active space and any document scope you add with @ tags.
- Uploads inherit the selected writable space.
- If expected documents are missing, first confirm that the active space is correct.

Ask Grounded Questions

1. Open Query Intelligence.
2. Confirm the active Knowledge Space.
3. Choose Auto, Documents, Live DB, or Hybrid. Select one approved DB source when needed.
4. Type a specific question. Mention dates, document names, entities, or constraints when useful.
5. Press Enter or click Send. Use Shift+Enter for a new line.
6. Watch the progress panel as the system routes the query, retrieves evidence, and drafts the answer.
7. Read the answer and inspect the citations.

Choose a Query Source

Mode	Use it for
Auto	Let Prudentia choose visible documents, approved live data, or both.
Documents	Use indexed files in the active space and optional @ document tags.
Live DB	Use current structured facts from approved SQL Server or PostgreSQL scopes.
Hybrid	Combine indexed document evidence with live database rows.

In Live DB or Hybrid mode, select one source or leave All visible DB sources selected. If a selected source cannot answer, use Search all sources when offered.

Verify Answers With Evidence

- PDF sources can show page navigation and highlighted regions.
- DOCX sources can show matching text where available.
- Image sources can show image highlights or extracted image evidence.
- Low faithfulness, degraded response, conflict, or grounding warnings mean you should verify manually before sharing or acting.

Use Chat History and Generated Files

- Completed conversations appear in the chat history rail.
- Use New chat to reset context.
- Reopen previous conversations when continuing the same task.
- Delete old conversations that are no longer useful.
- When document generation starts, answer clarification questions and download the DOCX, PDF, or PPTX when ready.

Use Optional GraphRAG

- Select GraphRAG during upload or reingestion only for documents that should participate in relationship and corpus-level analysis.
- Leaving GraphRAG off keeps standard document search available and does not create a graph task.
- Opted-in documents are searchable before the separate graph task necessarily finishes.
- Ask for themes, patterns, trends, recurring issues, risks, relationships, or executive overviews across opted-in documents.
- Graph answers cite source chunks. Exact counts and exhaustive lists use structured or standard retrieval.

Upload Documents

If Add Files appears in your sidebar, you can upload supported source files into a writable Knowledge Space.

Item	Details
Supported files	PDF, DOCX, JPG, JPEG, PNG, and JSON
Maximum size	50 MB per file
Required metadata	A writable Knowledge Space
Processing choice	Fast, Balanced, or High accuracy; optional GraphRAG opt-in
Optional metadata	Effective date, expiry date, description, and supersedes IDs for single-file replacements
Result	Each file receives its own ingestion job

Upload Workflow

1. Open Document Intake > Add Files.
2. Select one or more supported files.
3. Choose the writable Knowledge Space.
4. Choose the ingestion profile and enable GraphRAG only when graph analysis is needed.
5. Add dates, description, or supersession details if needed.
6. Click Upload documents.
7. Watch Recent upload jobs for progress and warnings.

Track Indexing Activity

Open Activity to monitor upload, folder, restore, and reingestion jobs.

Status	Meaning
Scheduled	Waiting for a planned start time.

Status	Meaning
Queued	Accepted and waiting for an ingestion worker.
Processing	Being parsed, enriched, chunked, embedded, and indexed.
Needs review	Paused for human OCR, extraction, or PDF image review.
Indexed	Complete and searchable.
Failed	Stopped with an error that needs investigation.
Cancelled	Manually stopped before completion.

- Filter by status, origin, Knowledge Space, date range, or search text.
- Open a job to review stage details, warnings, attempts, and parser provenance.
- If you have write access, you may cancel jobs that are still scheduled, queued, or processing.
- Use Ingestion Health to see active pipeline, failed jobs, review queues, and recent failures.
- GraphRAG activity shows queue, worker, and active-task status for documents that opted in to graph enrichment.

Browse the Document Library

- Document Overview shows active documents, current versions, processing items, attention items, and Trash count.
- Knowledge Spaces shows the folder tree and direct documents in selected spaces.
- Documents lets you search and filter by lifecycle and ingestion status.
- Trash contains soft-deleted documents that authorized users can restore or permanently delete.
- The Document Inspector shows metadata, extracted facts, version history, downloads, and allowed lifecycle actions.

Use Folder Sources

- Choose a folder from the browser to stage a point-in-time copy with relative paths preserved.
- Folder snapshots are one-time schedules. Re-browse the folder later to ingest newer local edits.
- Use schedule run details to review queued, skipped, and failed files.

Use Database Connectors

Platform Admins, System Admins, and Space Admins can prepare SQL Server or PostgreSQL for governed live read-only answers. Record-sync schedules are retired; rows are queried at question time and are not copied into the document corpus.

1. Add a read-only profile with host, port, database, credentials, and required driver or SSL mode.
2. Save the encrypted profile and click Test.
3. Click Read Schema to inspect tables, columns, relationships, indexes, and raw schema JSON.
4. Click Prepare Review and review names, business rules, descriptions, synonyms, allowed fields, and sensitivity flags.
5. Assign the owner Knowledge Space, optional shared Knowledge Spaces, and clearance, then Enable Live DB Access after review.
6. Choose Live DB or Hybrid in Query Intelligence and select the approved source.

Only approved, non-sensitive tables, columns, and relationships can be queried. Write statements, wildcards, unapproved joins, system schemas, and chained SQL are blocked.

Use Review Queue - OCR Blocks

If Review Queue appears in your sidebar, use it to resolve OCR blocks before indexing continues.

1. Open Review Queue and choose OCR blocks.
2. Select a pending block grouped under a document.
3. Compare the OCR text with the source page preview.
4. Edit Corrected extraction text.
5. Click Approve correction to save, or Reject block if it should not continue.

Use Review Queue - PDF Images

Large PDF image-analysis batches can pause here for human triage. The job resumes when no candidates remain pending.

1. Choose PDF images.
2. Select a held document image batch.
3. Review candidate thumbnails and recommended markers.
4. Choose Analyze selected, Analyze recommended, Skip selected, or Skip all pending.

Read the System Audit

If System Audit appears in your sidebar, use it to review visible authentication, document, ingestion, and workspace events.

- Events are newest first.
- Each row shows event type, target, actor, payload summary, and created time.
- Audit is for review and escalation. It is not an editing surface.

Use Runtime Settings

Runtime Settings is Platform Admin-only under Govern. It contains Models & Roles, Inference Services, and Ingestion Controls.

Models & Roles

- Choose an Ollama or vLLM starting stack, check service discovery, and assign synthesis, reasoning, router, faithfulness, ingestion, vision, embedding, and reranker roles.
- Embedding changes require document reindexing.
- Tune output and evidence budgets, timeouts, Ollama thinking, and the query planner.
- Test draft validates without activation; Save & activate validates again and makes the draft active without a container restart.

Inference Services

Check provider availability and manage vLLM text, embedding, and vision limits. Applying limits restarts only the selected service and interrupts requests that use it.

Ingestion Controls

Profile	Behavior	Best for
Fast	Native text first, no vision, and capped Docling repair.	High-volume, text-native files
Balanced	Native text first with a larger Docling repair budget.	Mixed collections and moderate layout complexity
High accuracy	Deeper layout repair with preference for full-document repair.	Scans, tables, and fidelity-sensitive layouts

- Worker concurrency is 1-10 per replica; one is recommended for the default 8 GB environment.
- OCR below the review threshold pauses in Review Queue.
- PDF image batches above the image review threshold pause before vision analysis; set the threshold to 0 to skip this gate.
- Vision layout repair re-reads complex PDF pages after Docling and trades throughput for fidelity.
- Graph enrichment controls whether completed documents expose the graph enrichment action.

Troubleshooting

Issue	Likely cause	What to do
A page is missing	Your account does not have that route.	Ask an admin to confirm your role and Knowledge Space memberships.
You cannot query	The active space has no indexed current documents, or the model runtime is unavailable.	Switch space, wait for indexing, or report the runtime issue.
No documents appear	Wrong active space, insufficient clearance, or no documents are current.	Check the space selector and ask your workspace admin to verify access.
Upload is rejected	Unsupported type, empty file, file over 50 MB, or no writable space.	Use PDF, DOCX, JPG, JPEG, PNG, or JSON and choose a writable Knowledge Space.
A document needs review	OCR, extraction, or PDF image triage paused indexing.	A Document Contributor, Space Admin, System Admin, or Platform Admin should resolve the pending Review Queue item.
An ingestion job looks stuck	Worker heartbeat, parsing, OCR, or model services may be delayed.	Open Activity or Ingestion Health and share the job ID with operations.
Source highlight is unavailable	The source may lack page coordinates or exact text matching.	Use the original document link and compare page, excerpt, and citation details.
The answer has a warning	Sources may conflict, faithfulness may be low, or the response is degraded.	Inspect citations manually before using the answer.
No Live DB source appears	No approved catalog is visible for your space and clearance.	Ask a connector admin to test the profile and approve the minimum required schema scope.
Graph enrichment is absent	GraphRAG was not selected, is disabled, or its worker is unavailable.	Standard retrieval still works. Check the GraphRAG choice and Ingestion Health when graph analysis is needed.

Safe Usage Practices

- Do not rely on an answer without checking citations for important decisions.
- Use the correct Knowledge Space before querying or uploading.
- Do not upload documents above your clearance or outside your team scope.
- When replacing a file, use supersedes metadata so the older version stays auditable.
- Do not share generated files beyond the same access and clearance rules as the source documents.
- Report repeated model, indexing, or source-viewer errors with the job ID, document title, and time observed.

Glossary

Term	Meaning
Knowledge Space	A document access scope such as /finance or /manuals.
Current document	The active version used for retrieval.
Superseded document	An older version kept for audit and history.
Trash	Soft-deleted documents that can be restored by authorized users.
Ingestion job	The background task that parses, enriches, embeds, and indexes a file.
Citation	A link between an answer claim and its supporting source evidence.
Faithfulness	A grounding check that estimates whether the answer is supported by retrieved evidence.
Generated artifact	A downloadable DOCX, PDF, or PPTX produced from a qualifying request.
Approved database scope	A reviewed set of database tables, columns, relationships, space, and clearance allowed for live read-only SQL.
Ingestion profile	Fast, Balanced, or High accuracy parsing selected for a new document.
GraphRAG	Optional entity-and-relationship enrichment for corpus-level themes and patterns.